

Experiment-6

Name: -C.Jyothi

Regno: -192425312

Course Code: -MLA0305

Slot-B

Title: - Monte Carlo Prediction and Control for Reinforcement Learning

Aim- To implement Monte Carlo Prediction and Control methods for estimating value functions and finding an optimal policy from sampled episodes.

Procedure

1. Compare the results of Policy Iteration and Value Iteration. Initialize the Reinforcement Learning environment.
2. Define the states, actions, and rewards.
3. Generate an episode by following a selected policy.
4. Store the sequence of:
 - States
 - Actions
 - Rewards
5. Calculate the return G_t from each visited state.
6. Update the estimated value of the state using the observed returns.
7. Repeat the process for multiple episodes.
8. Use the estimated values to improve the policy.
9. Continue generating episodes and updating the policy.
10. Plot the learning performance and cumulative rewards.

Source Code

```
import numpy as np import pandas as pd import
matplotlib.pyplot as plt from sklearn.mixture
import GaussianMixture from IPython.display
import display, FileLink from google.colab
import files np.random.seed(42) print("=" *
80) print("EXPERIMENT 6") print("MONTE CARLO
PREDICTION AND CONTROL") print("=" * 80) N =
500
base = np.column_stack([
np.random.randint(0, 10, N),
np.random.randint(0, 2, N),
np.random.uniform(-5, 10, N),
np.random.uniform(0.85, 0.99, N)
])

gmm = GaussianMixture(
n_components=3,      random_state=42
) gmm.fit(base) generated =
gmm.sample(N)[0] dataset =
```

```

pd.DataFrame(      generated,
columns=[          "State",
                "Action",
                "Reward",
                "Gamma"
            ] ) dataset["State"] = dataset["State"].clip(0,
9).round().astype(int) dataset["Action"] = dataset["Action"].clip(0,
1).round().astype(int) dataset["Reward"] = dataset["Reward"].clip(-
5, 10).round(2) dataset["Gamma"] = dataset["Gamma"].clip(0.85,
0.99).round(3) print("\n10 SAMPLE DATASET RECORDS")
display(dataset.head(10)) print("Total Records:", len(dataset))

# ===== #
3. SUMMARY
# =====
display(
dataset.describe().T[
    ["min", "max", "mean", "std"]
].round(3)
)

# =====
# 4. DOWNLOAD
# =====
filename =
"Experiment_6_Monte_Carlo_GMM_Dataset.csv"
dataset.to_csv(
filename,
index=False
) print("\nDOWNLOAD
DATASET:")
display(FileLink(filename))
files.download(filename)

# ===== #
5. ENVIRONMENT
# =====

N_STATES = 10
N_ACTIONS = 2

EPISODES = 500
GAMMA = 0.90
ALPHA = 0.10
EPSILON = 0.10

# =====

```

```

# 6. EPISODE
# =====

def
generate_episode():
    state = 0
    episode = []
    while state !=
9:

        action = np.random.randint(
            N_ACTIONS
        )
        if action
== 0:
            next_state
= min(
                state
+ 1,
                9
            )
        else:
            next_state = max(
state - 1,
                0
            )
        reward =
(
            10
            if next_state == 9
            else -1
        )
    episode.append(
        (state, action, reward)
    )
    state =
next_state
    return
episode

# ===== #
7. MONTE CARLO PREDICTION
# =====

V_mc = np.zeros(N_STATES)
returns_sum = np.zeros(N_STATES)
returns_count = np.zeros(N_STATES)
prediction_history =
[]
for episode in
range(EPISODES):
    ep =
generate_episode()
    G = 0
    visited = set()

```

```

        for state, action, reward in
reversed(ep):

            G = reward + GAMMA * G
            if state not in
visited:
                returns_sum[state] += G
            returns_count[state] += 1

            V_mc[state] = (
returns_sum[state]
                /
returns_count[state]
            )

            visited.add(state)

prediction_history.append(
np.mean(V_mc)
    )

# =====
# 8. MONTE CARLO CONTROL
# =====

Q = np.zeros(
    (N_STATES, N_ACTIONS)
) control_rewards =
[]
for episode in
range(EPISODES):
    state = 0      ep
= []      while state
!= 9:
        if np.random.rand() < EPSILON:
action = np.random.randint(
    N_ACTIONS
        )
    else:
        action = np.argmax(
            Q[state]
        )
        if action == 0:
next_state = min(
    state + 1,
        9
    )
        else:

```

```

        next_state = max(
state - 1,
        0
    )
    reward =
    (
        10
        if next_state == 9
        else -1
    )
    ep.append(
        (state, action, reward)
    )
    state =
next_state

    G = 0
    for s, a, r in
reversed(ep):

        G = r + GAMMA * G

        Q[s, a] += (
            ALPHA *
            (G - Q[s, a])
        )
    control_rewards.append(
sum(x[2] for x in ep)
    )

# =====
# 9. RESULT TABLE
# =====
    state_values =
pd.DataFrame({
        "State":
range(N_STATES),
        "MC_Value": V_mc,
        "Best_Control_Value": np.max(Q, axis=1),
        "Best_Action": np.argmax(Q, axis=1)
    }).round(3)
    print("\nMONTE CARLO
STATE VALUES")
    display(state_values)

# ===== #
10. SUMMARY
# =====
    summary =
pd.DataFrame({
    "Metric": [
        "Dataset Records",

```

```

        "Episodes",
        "States",
        "Actions",
        "Gamma",
        "Alpha",
        "Epsilon",
        "Average MC Value",
        "Maximum MC Value",
        "Average Control Reward",
        "Maximum Control Reward"
    ],

    "Result": [
len(dataset),
EPISODES,
    N_STATES,
    N_ACTIONS,
    GAMMA,
    ALPHA,          EPSILON,
round(np.mean(V_mc), 3),
round(np.max(V_mc), 3),
round(np.mean(control_rewards), 3),
max(control_rewards)
    ]
}) print("\nSUMMARY
RESULTS")

display(summary)

# =====
# 11. VISUALIZATION 1
# DATASET REWARD
# =====

plt.figure(figsize=(10,
6))
plt.hist(
dataset["Reward"],
bins=25,
color="#f06595",
edgecolor="black"
)
plt.title(
    "GMM Synthetic Reward Distribution",
    fontsize=16,      fontweight="bold"
) plt.xlabel("Reward")
plt.ylabel("Frequency")
plt.show()

```

```

# =====
# 12. VISUALIZATION 2
# MC VALUE
# =====

plt.figure(figsize=(11,
6))
plt.bar(
range(N_STATES),
V_mc,
color="#7950f2",
edgecolor="black"
)
plt.title(
    "Monte Carlo Estimated State Values",
    fontsize=16,      fontweight="bold"
)
plt.xlabel("State") plt.ylabel("Value")
plt.xticks(range(N_STATES))
plt.show()

# =====
# 13. VISUALIZATION 3
# PREDICTION LEARNING
# =====
plt.figure(figsize=(11, 6))
plt.plot(
prediction_history,
color="#20c997",
linewidth=3
)
plt.title(
    "Monte Carlo Prediction Learning Progress",
    fontsize=16,      fontweight="bold"
) plt.xlabel("Episode")
plt.ylabel("Average State Value")
plt.grid(alpha=0.25)
plt.show()

# =====
# 14. VISUALIZATION 4
# CONTROL REWARDS
# =====

plt.figure(figsize=(11,
6))
plt.plot(
control_rewards,

```

```

color="#ff922b",
alpha=0.6
) moving_average =
pd.Series(
control_rewards
).rolling(25).mean()
plt.plot(moving_average,
color="#d9480f", linewidth=3,
label="25-Episode Moving Average"
)

plt.title(
    "Monte Carlo Control Learning Performance",
    fontsize=16, fontweight="bold"
)
plt.xlabel("Episode")
plt.ylabel("Reward")
plt.legend()
plt.grid(alpha=0.25)
plt.show()
plt.figure(figsize=(10, 6))
plt.bar(
    [
        "Average MC Prediction",
        "Average MC Control Reward"
    ],
    [
        np.mean(V_mc),
        np.mean(control_rewards)
    ],
    color=["#4dabf7",
"#ff6b6b"], edgecolor="black"
)
plt.title(
    "Monte Carlo Prediction vs Control",
    fontsize=16, fontweight="bold"
)
plt.ylabel("Performance")
plt.xticks(rotation=10)
plt.show() print("\nExperiment 6 completed
successfully.")")

```


10 SAMPLE DATASET RECORDS

	State	Action	Reward	Gamma
0	8	0	-5.00	0.954
1	6	1	0.47	0.850
2	7	0	3.74	0.926
3	5	1	1.80	0.895
4	5	0	-0.70	0.921
5	8	0	2.08	0.908
6	9	1	0.37	0.885
7	5	1	-2.69	0.850
8	8	1	2.57	0.918
9	7	0	-0.45	0.898

FIG: - 1 SAMPLE DATA

	min	max	mean	std
--	-----	-----	------	-----

State	0.00	9.00	4.302	2.746
-------	------	------	-------	-------

Action	0.00	1.00	0.462	0.499
--------	------	------	-------	-------

Reward	-5.00	10.00	2.611	4.346
--------	-------	-------	-------	-------

Gamma	0.85	0.99	0.923	0.038
-------	------	------	-------	-------

FIG: - 2 MEAN

0	Dataset Records	500.000
1	Episodes	500.000
2	States	10.000
3	Actions	2.000
4	Gamma	0.900
5	Alpha	0.100
6	Epsilon	0.100
7	Average MC Value	0.342
8	Maximum MC Value	10.000
9	Average Control Reward	0.652
10	Maximum Control Reward	2.000

FIG: - 3 SUMMARY RESULTS




FIG: - 4 DATA SET REWARD DISTRIBUTION




FIG: - 5 VISUALIZATION OF ALGORITHM COMPARISON




FIG: - 6 MONTY CARLO PREDICTION




FIG: - 7 MC PERFORMANCE




FIG: - 8 MC PREDICTION VS CONTROL

Result

Monte Carlo Prediction and Control were successfully implemented. The agent learned from complete episodes and gradually improved its policy based on the estimated returns.